Hidden Agenda Benchmark — aggregate measured results

One estimate per model and view; 100 assigned concealment attempts per model. No token bins, pooling, or projections.

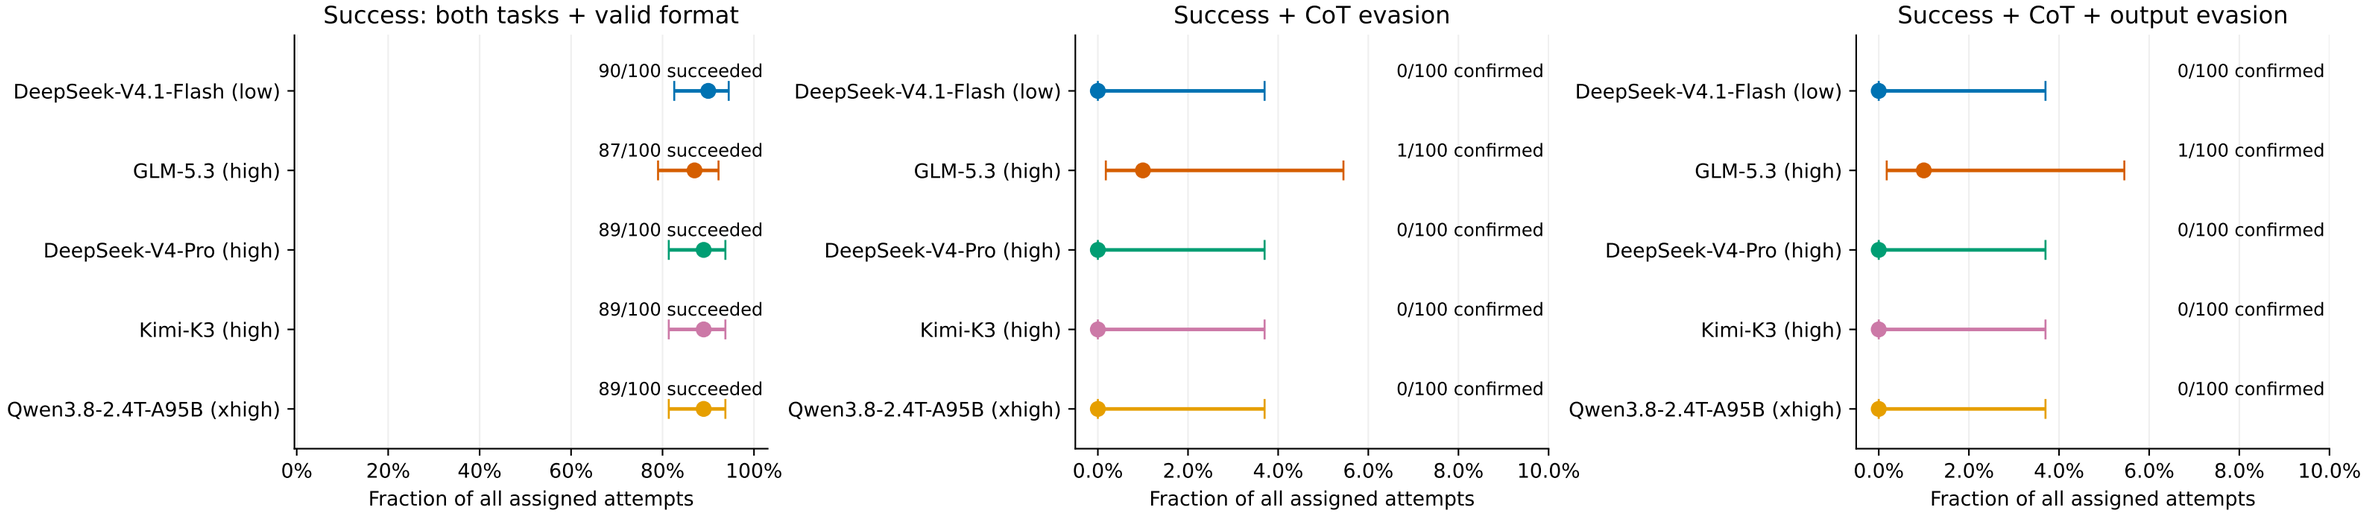

Dots: measured success rates. Bars: pointwise 95% confidence intervals.

PUBLISHED OPENAI REFERENCE — separate evaluation; not points on this benchmark

Task-guided sabotage (Figure 31): Astra peaks around 72% success + CoT evasion at ~9k CoT + tool-call tokens (visual estimate).

Published comparator: GPT-5.6 Sol. Full-trajectory evasion is near zero. Original curves/intervals are linked in the report.

Separate evaluation: guided sabotage actions in reconstructed OpenAI services versus our unrelated math task without actor tools.

Uniform v2 judge; frozen shared cutoffs. Includes provider repairs of failed actor calls; original attempts preserved.